

Monitoring Databricks clusters performance with Log Analytics - KSQL

Last updated by | Romanth Nirmal | Jun 24, 2019 at 4:52 PM GMT+1

Failed tasks

SparkListenerEvent_CL

```
| project TimeGenerated, Task_Info_Failed_b, 1  
| where isempty(Task_Info_Failed_b) == "true"  
| where TimeGenerated > ago(1d)
```

Driver memory used

SparkMetric_CL

```
| project name_s, TimeGenerated, value_d  
| where name_s contains "driver.jvm.total.used"
```

Shuffle traffic

SparkMetric_CL

```
| project name_s, TimeGenerated, value_d  
| where name_s contains "ExternalShuffle.shuffle-client.usedDirectMemory"
```

Count of total jobs completed - Driver level based on DAG

SparkMetric_CL

```
| project name_s, TimeGenerated, value_d  
| where name_s contains "driver.DAGScheduler.job.allJobs"
```

Count of Active jobs

SparkMetric_CL

| project name_s,TimeGenerated, value_d

| where name_s contains "DAGScheduler.job.activeJobs"

Count of total jobs completed - Worker level

SparkMetric_CL

| project name_s,TimeGenerated, value_d

| where name_s contains "executor.threadpool.completeTasks"

Example monitoring dashboard

<https://portal.azure.com/#@anglianwater.co.uk/dashboard/arm/subscriptions/8db87f54-8a28-40ff-b3f6-94aa11911faf/resourcegroups/aws-npe-dev-awbi-common-rg-n-eu/providers/microsoft.portal/dashboards/5dd49c7c-43c3-425d-ba99-05246bd2f055> 

